

LMJuly18-06 — LaunchMind: Build State & Roadmap

Date: July 18, 2026 · Part 6 of 6

Table of Contents

- [Current Branch & Git State](#)
 - [Milestone Summary](#)
 - [Test Suite State](#)
 - [TypeScript Status](#)
 - [UX Remediation v1.0 \(July 12\)](#)
 - [Design System v1.0 \(July 12\)](#)
 - [Morning Brief Refinements \(July 16–18\)](#)
 - [Pending Work](#)
 - [Pre-Launch Ops Checklist](#)
 - [65 ADRs Index](#)
 - [Seed Data State](#)
 - [Key Architectural Decisions](#)
-

1. Current Branch & Git State

Active branch: `july6addnewarchsec` **Hosted Supabase** (`gseqtbwdenjkwysregpp`): Migrations pushed through M03 (migrations 026–030). ⚠️ **BLOCKING:** Migrations 031–061 must be pushed before production traffic. See §9.
Single env file: `.env.local` (gitignored) — loads all 35 keys for local dev.

2. Milestone Summary

All 12 milestones are COMPLETE as of July 10, 2026.

Milestone 01 — Foundation (July 8, 2026)

- Architecture Review · 8 ADRs (ADR-001–008) · Architecture Baseline v1.0 · Blueprint v2.0
- Design System tokens · `PageShell`, `MetricCard`, `MissionCard`, `OpportunityCard`, `ApprovalCard`, `EmptyState`, `LoadingState` components

- Sidebar refactor (Tabler icons v3) · 15 new routes · Intelligence layout group
- Backend response envelope: `backend/src/lib/response.ts`, `ok()`, `fail()`, `ErrorCodes`
- URL redirects: `/briefs→/content`, `/insights→/results`, `/workspaces→/settings`

Milestone 02 — Product Workspace & Product Intake (July 8, 2026)

- 4 ADRs (ADR-011–014) · 3 migrations (032–034)
- `workspaceService.ts` + `integrationService.ts`
- Workspaces route extended (member management, workspace activate, product activate)
- Products route extended (Intake V3: setup/start, intake/step/:step, intake/complete, intake/status)
- Channels route (GA4, Firebase, website integrations)
- 5-step intake wizard: `/dashboard/products/setup/` → basics → business → audience → brand → connect

Milestone 03 — Content OS (Weeks 18–20)

- 3 migrations (026–028): `content_assets`, `content_preferences`, `learning_loop`
- `lib/aiClient.ts` — `callSonnet()` + `callHaiku()` with lazy Anthropic init
- `lib/creatomateClient.ts` — video render (graceful mock if key missing)
- `lib/elevenLabsClient.ts` — voice synthesis (graceful mock if key missing)
- `services/contentService.ts` — 6-step pipeline: context → Sonnet → Haiku scoring → ElevenLabs → Creatomate → DB
- `routes/contentAssets.route.ts` + `routes/settings.route.ts`
- Marketing image pipeline: `lib/replicateClient.ts` (Flux.1 Schnell) + `services/marketingImagesService.ts`
- Image decision tree: real screenshot (mockup) / Flux.1 (photo/graphic)
- Migration 029: product archive (4 archive routes)

Milestone 04 — Marketing Memory & Knowledge Graph (July 8, 2026)

- 4 ADRs (ADR-019–022) · 6 migrations (035–040)
- `marketing_memories`, `marketing_memory_versions`, `knowledge_nodes`, `knowledge_edges`, `evidence`, `learning_events`
- 3 backend services: `marketingMemoryService`, `knowledgeGraphService`, `learningPipelineService`
- 2 backend routes: `memory.route.ts` (9 routes) + `knowledge.route.ts` (7 routes)
- 2 frontend pages: Memory dashboard + Knowledge Graph explorer
- 17 tests passing

Milestone 05 — Context Engine & AI Platform (July 8, 2026)

- 5 ADRs (ADR-023–027) · 3 migrations (041–043)
- `lib/contextEngine.ts` (6 parallel sources, non-fatal) + `lib/promptRegistry.ts` + `lib/modelRouter.ts`
- `lib/aiPlatform.ts` — mandatory AI entry point: `callSonnet`, `callHaiku`, `callMessages`, `generateAI`

- Retry (2×), timeout (60s/30s), prompt injection defense, cost tracking
- 11 seeded prompts (migration 043)
- ALL services migrated to aiPlatform (no direct SDK calls outside aiClient.ts)
- `routes/ai.route.ts` (6 routes) + AI Audit frontend page
- 25 tests passing

Milestone 06 — Agent Platform & Mission Orchestrator (July 8, 2026)

- 5 ADRs (ADR-028–032) · 2 migrations (044–045)
- `backend/src/types/mission.ts` — all enums, DB interfaces, Zod schemas
- `services/missionService.ts` — full lifecycle management (15 methods)
- 12 agents (6 full, 6 stubs) + `agentRegistry.ts` dispatch table
- `workers/missionWorker.ts` — BullMQ queue, concurrency=5, DLQ via DB
- `routes/missions.route.ts` (9 routes)
- 2 frontend pages: Mission Center + Mission Detail (5s auto-poll when running)
- 17 tests passing

Milestone 07 — Owner Experience (July 8, 2026)

- 6 ADRs (ADR-033–038) · 1 migration (046): `saved_opportunities` + `notifications`
- `backend/src/routes/owner.route.ts` (9 adapter endpoints)
- Adapter pattern: `/owner/*` aggregates from existing services
- `/dashboard` redirects to `/dashboard/brief`
- 8 frontend pages: Morning Brief, Opportunities, Ask, Approvals, Results, Timeline, Ideas, Growth Brain
- `api.owner` namespace (8 methods) + related types
- 19 tests passing

Milestone 08 — Content Studio (July 8, 2026)

- 4 ADRs (ADR-039–042) · 4 migrations (047–050)
- `content_versions` (append-only) + `asset_approvals` + `publishing_targets` + `content_assets_extend`
- `routes/studio.route.ts` (9 endpoints) — generate, list, get, update, transform, versions, archive, restore, publish
- 5 new asset types: `blog_post`, `landing_page_copy`, `push_notification`, `release_notes`, `press_release`
- \$1.5 approval gate enforced: `POST /publish → 422 if approved_at is null`
- Version snapshot created BEFORE every update
- Full Content Studio page (Library / Generate / Stats tabs)
- 22 tests passing

Milestone 09 — Campaigns, Experiments & Execution (July 8, 2026)

- 7 ADRs (ADR-043–049) · 5 migrations (051–055)
- `campaigns_extend` + `experiments` + `experiment_variants` + `campaign_approvals` + `campaign_publish_attempts` + `execution_calendar_events`
- 3 backend routes: `campaigns.route.ts`, `experiments.route.ts`, `calendar.route.ts`
- A/B experiment framework with winner selection → learning pipeline
- Calendar: month view + list view, merges authored + derived events at API layer
- §1.5 + §1.6 enforced in campaign launch/schedule
- `api.campaigns`, `api.experiments`, `api.calendar` namespaces
- 23 tests passing (10 campaigns + 13 experiments)

Milestone 10 — Intelligence Network & Recommendation Engine (July 9, 2026)

- 4 ADRs (ADR-050–053) · 4 migrations (056–059)
- `decision_rules` (8 seeded) + `recommendation_feedback` + `intelligence_trends` + `saved_opportunities_m10` extension
- `decisionEngineService.ts` — 8 pure-TS rules, zero AI calls, AI cannot override
- `intelligenceNetworkService.ts` — min cohort=3 privacy guard, anonymous aggregation
- `recommendationEngineService.ts` — scoring formula ($\text{impact} \times 0.4 + \text{confidence} \times 0.3 + \text{urgency} \times 0.2 + \text{source} \times 0.1$)
- Market Intelligence page + Reviews Intelligence page (both with real scraped_meta data)
- 28 tests passing (including tenant isolation verification)

Milestone 11 — Analytics, Reporting & Optimization (July 9, 2026)

- 4 ADRs (ADR-054–057) · 2 migrations (060–061)
- `reports` (AI narrative cache) + `optimization_insights`
- `analyticsService.ts` (summary, KPI trend, attribution, funnel, ROI)
- `reportingService.ts` (5 report types, cache-first, triggers learning pipeline)
- `optimizationEngineService.ts` (6 insight types, high-confidence → auto-recommendations)
- `routes/analytics.route.ts` (8 endpoints) + `routes/reports.route.ts` (5 endpoints)
- Analytics page (KPI cards, funnel, ROI table, insights) + Reports page (generate, drawer, export)
- 29 tests passing (16 analytics + 13 reports)

Milestone 12 — Production Hardening & Enterprise Readiness (July 10, 2026)

- 8 ADRs (ADR-058–065) · 10 review documents · Final readiness report
- 0 CRITICAL, 0 HIGH security findings
- 349/351 tests passing (2 pre-existing non-blocking failures)
- VERDICT: **APPROVED FOR PRODUCTION** pending 10 pre-launch ops tasks

3. Test Suite State

Last run: 349/351 passing

Test file	Tests	Status
auth.test.ts	~12	✓
products.test.ts	~18	✓
channels.test.ts	~8	✓
campaigns.test.ts	10	✓
experiments.test.ts	13	✓
briefs.test.ts	~10	✓
founders.test.ts	~12	✓
workspaces.test.ts	~10	✓
memory.test.ts	17	✓
aiPlatform.test.ts	25	⚠ 1 pre-existing failure (fake timers)
missions.test.ts	17	✓
owner.test.ts	19	✓
studio.test.ts	22	✓
recommendations.test.ts	28	✓
analytics.test.ts	16	✓
reports.test.ts	13	✓
content.test.ts	~5	⚠ 1 pre-existing failure (mock shape)

2 pre-existing failures (non-blocking, unchanged since M08):

- content.test.ts — mock shape mismatch in one test (doesn't affect runtime)
- aiPlatform.test.ts — fake timer conflict in one test (doesn't affect runtime)

4. TypeScript Status

Frontend (npx tsc --noEmit): 0 errors Backend (cd backend && npx tsc --noEmit): 5 pre-existing errors in scraper layer

- `scraperQueue.ts` — BullMQ type drift (library update changed generic shape)
- `icpService.ts` — Playwright type drift (library update changed `page.goto` return type)
- `scraperWorker.ts` — same Playwright drift
- These have no runtime impact. Fix in a dedicated `chore(types): fix scraper library type drift` commit.
- All M10–M12 files: 0 errors

5. UX Remediation v1.0 (July 12, 2026)

R1 — Evidence Crash Fix (CRITICAL)

- `Opportunity.evidence` typed as `unknown` in `lib/api.ts` (was `string[] | null` — incorrect)
- `opportunities/page.tsx`, `brief/page.tsx`, `ask/page.tsx` — use `toStringArray()` at call sites
- `EvidenceChips` props widened to `chips: unknown`; internal `toStringArray()`
- `owner.route.ts` — evidence normalized to `array[]` on GET /owner/brief + GET /owner/opportunities
- `lib/coerce.ts` — `toStringArray()` + `toRecord()` defensive JSONB coercion
- `lib/__tests__/coerce.test.ts` — 14 unit tests

R2 — JSONB Audit

- `market/page.tsx` — `competitor_set` object shape handled
- `reviews/page.tsx` — `scraped_meta` reads via `toRecord()`, reviews + themes via `Array.isArray`

R3 — Morning Brief Progressive Render

- `RecommendationSkeleton` — pulsing skeleton while AI rec loads
- `RecommendationUnavailable` — friendly fallback with "Try again" button
- `BriefPage` tracks `recState: 'loading' 'ready' 'failed'`
- `loadBrief()` is `useCallback` returning cleanup; 8s hard ceiling timer; retry without page reload

R4 — ErrorState + Hardened Error Views

- `components/launchmind/ErrorState.tsx` — third member of state trio (Loading / Empty / Error)
- `market/page.tsx` + `reviews/page.tsx` + `analytics/page.tsx` — raw error strings replaced with `<ErrorState onRetry>`

R5 — Responsive Mobile Nav

- `components/launchmind/MobileNav.tsx` — 5-item bottom tab bar, `lg:hidden`, sidebar-dark bg, iOS safe-area-inset-bottom

- `app/(dashboard)/layout.tsx` — renders `<MobileNav>` , main gets `pb-16 lg:pb-0`
- `Sidebar.tsx` — `hidden lg:flex` (was always visible regardless of viewport)

R6 — Growth Brain Product Picker

- `growth-brain/page.tsx` — `allProducts` state; distinguishes "no products" vs "no active product"; multi-product switcher shown when >1 product

6. Design System v1.0 (July 12, 2026)

Basis: `LaunchMind-Design-System-v1.0.md` — authoritative spec for all future UI.

Token Changes

- Removed shadcn/Radix HSL var block (shadcn not installed — §16 Option A)
- `--red` / `--red-d` / `--red-b` → renamed to `--danger` / `--danger-d` / `--danger-b`
- `--sidebar2` → `--sidebar-2` (hyphen convention)
- Added `--ai` / `--ai-d` / `--ai-b` / `--ai-l` (violet #7c5cff) — AI provenance ONLY
- Added `--r-full` , `--e1/--e2/--e3` (elevation), motion tokens
- Added `@media (prefers-reduced-motion)` block (mandatory)

Violet usage rule (strictly enforced): `var(--ai)` is for AI provenance only.

-  Permitted: `AIBadge` component, `ConfidenceBadge`, "What I learned" border-left accent, `WhyThisPanel`
-  Forbidden: buttons, backgrounds, gradients, non-AI content

New Components (M12 / Design System)

Component	Purpose
<code>AIBadge.tsx</code>	"♦ AI generated" violet badge (§10.1)
<code>ConfidenceBadge.tsx</code>	0–100 normalized confidence display (§10.2)
<code>EvidenceChips.tsx</code>	Defensive chip list from unknown JSONB data (§10.3)
<code>WhyThisPanel.tsx</code>	Expandable Why/Evidence/Confidence/Risk/Source panel (§10.4)
<code>Button.tsx</code>	Canonical primitive, 4 variants × 3 sizes (§11.3)
<code>ErrorState.tsx</code>	Error fallback with optional retry (from UX Remediation R4)
<code>MobileNav.tsx</code>	Bottom tab bar for mobile (from UX Remediation R5)

Token Rename — Global Replace (52 files)

- `var(--red)` → `var(--danger)` , `var(--red-d)` → `var(--danger-d)` , `var(--red-b)` → `var(--danger-b)`
- `sidebar2` → `sidebar-2` (CSS), `sidebarHover` in `lib/design-system/tokens.ts`

7. Morning Brief Refinements (July 16–18, 2026)

Four targeted improvements to `app/(dashboard)/dashboard/brief/page.tsx`:

Item 1 — Header Uses `recommendation.title`

Before: "Here's where {productName} stands today." / rec.summary duplicated below After: "I reviewed {productName} overnight. {rec.title}" — eliminates duplication with the RecommendationCard content below

Item 2 — 3-Up MetricCard Grid

Replaced 12px gray metadata line with three MetricCard components:

- **Installs this week** — `metrics.weeklyInstalls` + `weekOverWeekInstallDelta` as delta
- **Avg CPI** — `metrics.cpi` formatted as currency string
- **Active campaigns** — `metrics.activeCampaigns` count

3 columns on desktop (`sm:grid-cols-3`), stacked on mobile. Null = "—". No new fetches.

Item 3 — SinceThenStrip Deduplication

If the only bullet in SinceThenStrip is the approval count AND the amber approval banner is already visible (`pendingApprovals.total > 0`), return `null` . Prevents the same info showing twice.

Item 4 — "What I Learned" Panel

Backend addition to `owner.route.ts` : 7th parallel query fetches top-3 marketing_memories (by confidence, archived=false). Added `memories: Array<{id, title, body, memoryType, confidence}>` to `BriefResponse` in `lib/api.ts` .

Frontend panel (right column, below Growth Brain card):

- Only renders when `data.memories.length > 0`
- Left border: `2px solid var(--ai-b)` — violet only here, never fill
- Shows max 2 memories: `memoryType` label, `body ?? title` text, `ConfidenceBadge`
- "View memory →" link to `/dashboard/intelligence/memory`

8. Pending Work

Immediate (pre-launch blockers)

- Push migrations 035–061 to hosted Supabase (`gseqtbwdenjkwysregpp`)
- Set `ELEVENLABS_API_KEY` , `CREATOMATE_API_KEY` , `REPLICATE_API_TOKEN` on Oracle VM
- Create migration `062_production_indexes.sql` (covering indexes for hot query paths)
- Enable pgBouncer in Supabase + promote Upstash to paid plan (before 100+ founders)
- Publish privacy notice + designate India Grievance Officer (DPDP compliance)

Short-term (M13 scope)

- Fix 2 pre-existing test failures:
 - `content.test.ts` — mock shape mismatch - `aiPlatform.test.ts` — fake timer conflict
- Fix 5 pre-existing TypeScript errors (scraper layer library type drift)
- Implement 6 stub agents: `planningAgent` , `creativeAgent` , `publishingAgent` , `optimizationAgent` , `learningAgent` , `benchmarkAgent`
- Billing plan-change flow (requires live `STRIPE_SECRET_KEY` on Oracle VM)

Medium-term (M13 features)

- Actual platform API posting via channel adapters (Meta, Google, LinkedIn, WhatsApp Business)
- Requires: OAuth tokens + per-platform SDK integration
- Calendar: drag-and-drop reschedule (deferred — requires react-dnd or dnd-kit)
 - Add OpenTelemetry spans (Axiom logs cover observability for now)
 - Add `--coverage` flag to CI and enforce 80% gate on new files
 - Add SAST semgrep rule blocking direct Anthropic SDK use outside `aiClient.ts`
 - Growth Brain (Milestone 03): placeholder stub — Mission Orchestrator is now the execution layer

Deferred

- India market: PDPB 2023 full compliance review before enabling INR payments in production
 - Multi-workspace: tenant isolation for Studio plan (`workspace_members` table exists, but multi-tenant UX needs more work)
 - DAST: full OWASP ZAP run against staging before production traffic
-

9. Pre-Launch Ops Checklist

From docs/release/production-readiness-checklist.md :

#	Task	Status
1	Push migrations 035–061 to hosted Supabase	Pending
2	Verify all 61 migrations applied + RLS enabled	Pending
3	Set ELEVENLABS_API_KEY on Oracle VM	Pending
4	Set CREATOMATE_API_KEY on Oracle VM	Pending
5	Set REPLICATE_API_TOKEN on Oracle VM	Pending
6	Create + apply migration 062 (production indexes)	Pending
7	Enable pgBouncer connection pooling in Supabase	Pending
8	Promote Upstash Redis to paid plan	Pending
9	Publish privacy notice at launchmind.com/privacy	Pending
10	Designate India DPDP Grievance Officer	Pending
11	Run OWASP ZAP against staging	Pending
12	Verify SSRF IP blocklist in scraper worker	Pending
13	Enable Sentry alerting for ERROR+ level	Pending
14	Configure Axiom dashboards + alerts	Pending
15	Load test: 50 concurrent founders on strategy generation	Pending
16	Validate Stripe webhooks in production	Pending
17	Validate Razorpay webhooks in production	Pending
18	Verify Resend domain authentication	Pending
19	Enable Cloudflare WAF production rules	Pending
20	Document on-call runbook in docs/incidents/playbook.md	Pending

10. 65 ADRs Index

ADR	Title	Milestone
ADR-001	Fastify over Express	M01 Foundation

ADR-002	Supabase Auth over custom JWT	M01 Foundation
ADR-003	BullMQ over direct cron	M01 Foundation
ADR-004	pgvector over Pinecone	M01 Foundation
ADR-005	Additive-only migrations	M01 Foundation
ADR-006	Vitest over Jest	M01 Foundation
ADR-007	Tailwind + shadcn over custom CSS	M01 Foundation
ADR-008	URL redirect strategy (/briefs→/content etc.)	M01 Foundation
ADR-009	(reserved)	—
ADR-010	(reserved)	—
ADR-011	Workspace isolation model	M02 Workspaces
ADR-012	Intake V3 step schema	M02 Product Intake
ADR-013	Integration adapter pattern	M02 Integrations
ADR-014	Brand voice extraction strategy	M02 Brand
ADR-015	(reserved)	—
ADR-016	(reserved)	—
ADR-017	(reserved)	—
ADR-018	(reserved)	—
ADR-019	Marketing memory append-only versioning	M04 Memory
ADR-020	Postgres adjacency list for knowledge graph	M04 Knowledge
ADR-021	Single <code>ingestLearningEvent()</code> entry point	M04 Learning
ADR-022	Manual-only memory merge	M04 Memory
ADR-023	Context Engine architecture	M05 AI Platform
ADR-024	AI Platform mandatory entry point	M05 AI Platform
ADR-025	Prompt registry and versioning	M05 AI Platform
ADR-026	Model routing table	M05 AI Platform
ADR-027	AI request audit strategy	M05 AI Platform
ADR-028	Agent platform design	M06 Agents
ADR-029	Mission orchestrator lifecycle	M06 Agents
ADR-030	BullMQ queue strategy for missions	M06 Agents
ADR-031	Mission lifecycle state machine	M06 Agents

ADR-032	Agent isolation model	M06 Agents
ADR-033	Owner experience architecture	M07 Owner
ADR-034	Morning Brief replaces dashboard	M07 Owner
ADR-035	Ask LaunchMind as command center	M07 Owner
ADR-036	Opportunities as growth backlog	M07 Owner
ADR-037	Progressive disclosure UX	M07 Owner
ADR-038	Approval UX enforcement	M07 Owner
ADR-039	Unified content pipeline	M08 Studio
ADR-040	Asset library architecture	M08 Studio
ADR-041	Content versioning (append-only)	M08 Studio
ADR-042	Media integration (Creatomate / ElevenLabs)	M08 Studio
ADR-043	Campaign execution model	M09 Campaigns
ADR-044	Experiment framework	M09 Experiments
ADR-045	Channel adapter architecture	M09 Campaigns
ADR-046	Server-side approval enforcement	M09 Campaigns
ADR-047	Budget guardrails	M09 Campaigns
ADR-048	Publishing retry strategy	M09 Campaigns
ADR-049	Execution calendar	M09 Calendar
ADR-050	Intelligence network architecture	M10 Intelligence
ADR-051	Recommendation engine scoring	M10 Recommendations
ADR-052	Decision engine (pure TS, no AI)	M10 Decision
ADR-053	Anonymous benchmarking	M10 Intelligence
ADR-054	Unified analytics model	M11 Analytics
ADR-055	Reporting framework (cache-first)	M11 Reports
ADR-056	Optimization engine	M11 Optimization
ADR-057	Attribution strategy (last-touch)	M11 Analytics
ADR-058	Production security architecture	M12 Hardening
ADR-059	Compliance strategy (GDPR/CCPA/DPDP)	M12 Hardening
ADR-060	Performance and scalability	M12 Hardening
ADR-061	Observability and alerting	M12 Hardening

ADR-062	CI/CD deployment pipeline	M12 Hardening
ADR-063	AI safety and cost controls	M12 Hardening
ADR-064	Data protection and retention	M12 Hardening
ADR-065	Quality gate strategy	M12 Hardening

11. Seed Data State

All seed data is in hosted Supabase (`gseqtbwdenjkwysregpp`).

Founder

- **Email:** vijay@lm.com
- **Plan:** Solo (\$19/month)
- **Token balance:** 300
- **Onboarding step:** 6 (fully onboarded)

Product: ClientPulse

- **Store URL:** App Store
- **Platform:** app_store
- **intake_step:** 6 (intake complete)
- **Primary channel:** whatsapp
- **Primary market:** india
- **founder_context:** full 5-question responses
- **MOAT:** present in strategy

Campaigns (3 launched)

- WhatsApp / India — launched
- Meta / USA — launched
- Google / India — launched

Campaign Metrics

- 10 weeks of data for each campaign
- Mix of impressions, clicks, installs, CPI, CTR, ROAS

Weekly Brief

- 1 brief created, status=sent

Marketing Memories (5 rows)

Title	Type	Confidence
brand memory: clientpulse	brand	0.81
product memory: clientpulse	product	0.79
customer memory: clientpulse	customer	0.86
campaign memory: clientpulse	campaign	0.58
founder memory: clientpulse	founder	0.63

Playbook Signals

52 rows: 28 from migration 011 + 24 from migration 018

- Mix of USA + India markets
- All 6 channels represented
- Various price tiers and hook types

AI Requests

Populated after any strategy generation with vijay@lm.com.

12. Key Architectural Decisions

These are the decisions most likely to affect future development:

The AI Platform is inviolable

All AI calls MUST flow through `aiPlatform.ts`. There is a SAST rule blocking direct SDK use outside `aiClient.ts` (planned for M13). Every call is audited in `ai_requests`. The `auditCtx` object is required on every call — never optional.

Migrations are additive-only, forever

Migration numbers are permanent. If a column needs to be removed, it stays in the schema and is ignored by the application. If a type needs to change, add a new column and migrate data. This is non-negotiable.

The Decision Engine never calls AI

`decisionEngineService.ts` is pure TypeScript. It does zero AI calls. The 8 rule functions (§1.5 approve-before-post, §1.6 spend cap, plan gates, token balance, regen limit, experiment runtime, workspace permission, benchmark access) cannot be overridden by any AI recommendation or mission agent.

Content Studio versioning is append-only

`content_versions` has `REVOKE UPDATE, DELETE`. `asset_approvals` has `REVOKE UPDATE, DELETE`. Version history is permanent. Archive = soft delete (timestamp), never hard delete.

Intelligence Network is fully anonymous

`intelligence_trends` contains NO `founder_id`, NO `product_id`, NO `product_name`. Minimum cohort of 3 signals required before benchmark is published. Authenticated SELECT only; INSERT via `service_role` only.

The Context Engine is best-effort

6 parallel queries in `buildContextPackage()`. Any individual source failing (no analytics data, empty graph) is caught and treated as `undefined`. AI prompts must handle missing context gracefully. A failing source never aborts the full context build.

JWT verification is algorithm-agnostic

Supabase rotated from HS256 to ES256 on 2026-05-16. The backend uses `supabase.auth.getUser()` (not manual JWT decode) — algorithm-agnostic and future-proof against further Supabase key rotations.

Frontend never calls Supabase directly

The frontend calls the Fastify backend API (`NEXT_PUBLIC_API_URL`). RLS is a defence-in-depth measure, not the primary access control. All data access is through typed Fastify routes with Zod validation. `lib/api.ts` is the only API client file.

Morning Brief is the primary dashboard

`/dashboard` redirects to `/dashboard/brief`. The Morning Brief is the single entry point for the owner's daily workflow. The old `/dashboard` page (`app/(dashboard)/dashboard/page.tsx`) contains only a redirect component.

End of documentation. Files in this set:

- [LMJuly18-00-INDEX.md](#) — Master index
- [LMJuly18-01-Overview-Architecture.md](#) — Product, stack, rules, infrastructure
- [LMJuly18-02-Database-Schema.md](#) — 61 migrations, all table schemas, RLS

- [LMJuly18-03-Backend.md](#) — Routes, services, workers, lib, types
- [LMJuly18-04-Frontend.md](#) — Pages, components, design system, api.ts
- [LMJuly18-05-Intelligence-Agents.md](#) — AI platform, agents, intelligence systems
- **LMJuly18-06-Build-State-Roadmap.md** — This file: milestones, test state, pending work, ADRs